

**(a) Leaderboard:  
one number per method**

**BARO > CIRCA**

RCAEval RE1

**+0.147**

**RCD > BARO**


PetShop

**+0.135**


*pooled difference,  
k subsystems averaged*

**(b) Premise 1:  
on which systems does it hold?**

BARO vs CIRCA, RCAEval RE1  
same sign on all three applications  
margin 0.147 > spread 0.084



RCD vs BARO, PetShop  
**reversed on High traffic, CI excludes 0**  
margin 0.135 < spread 0.219



**(c) Premise 2:  
how was the score produced?**

CIRCA vs RCD, RCAEval RE1  
same cases, two input files

