

**(a) Leaderboard:
one number per method**

BARO > CIRCA

RCAEval RE1

+0.147

RCD > BARO


PetShop

+0.135


*pooled difference,
k subsystems averaged*

**(b) Premise 1:
on which systems does it hold?**

BARO vs CIRCA, RCAEval RE1
same sign on all three applications
margin 0.147 > spread 0.084



RCD vs BARO, PetShop
reversed on High traffic, CI excludes 0
margin 0.135 < spread 0.219



**(c) Premise 2:
how was the score produced?**

CIRCA vs RCD, RCAEval RE1
same cases, two input files

